

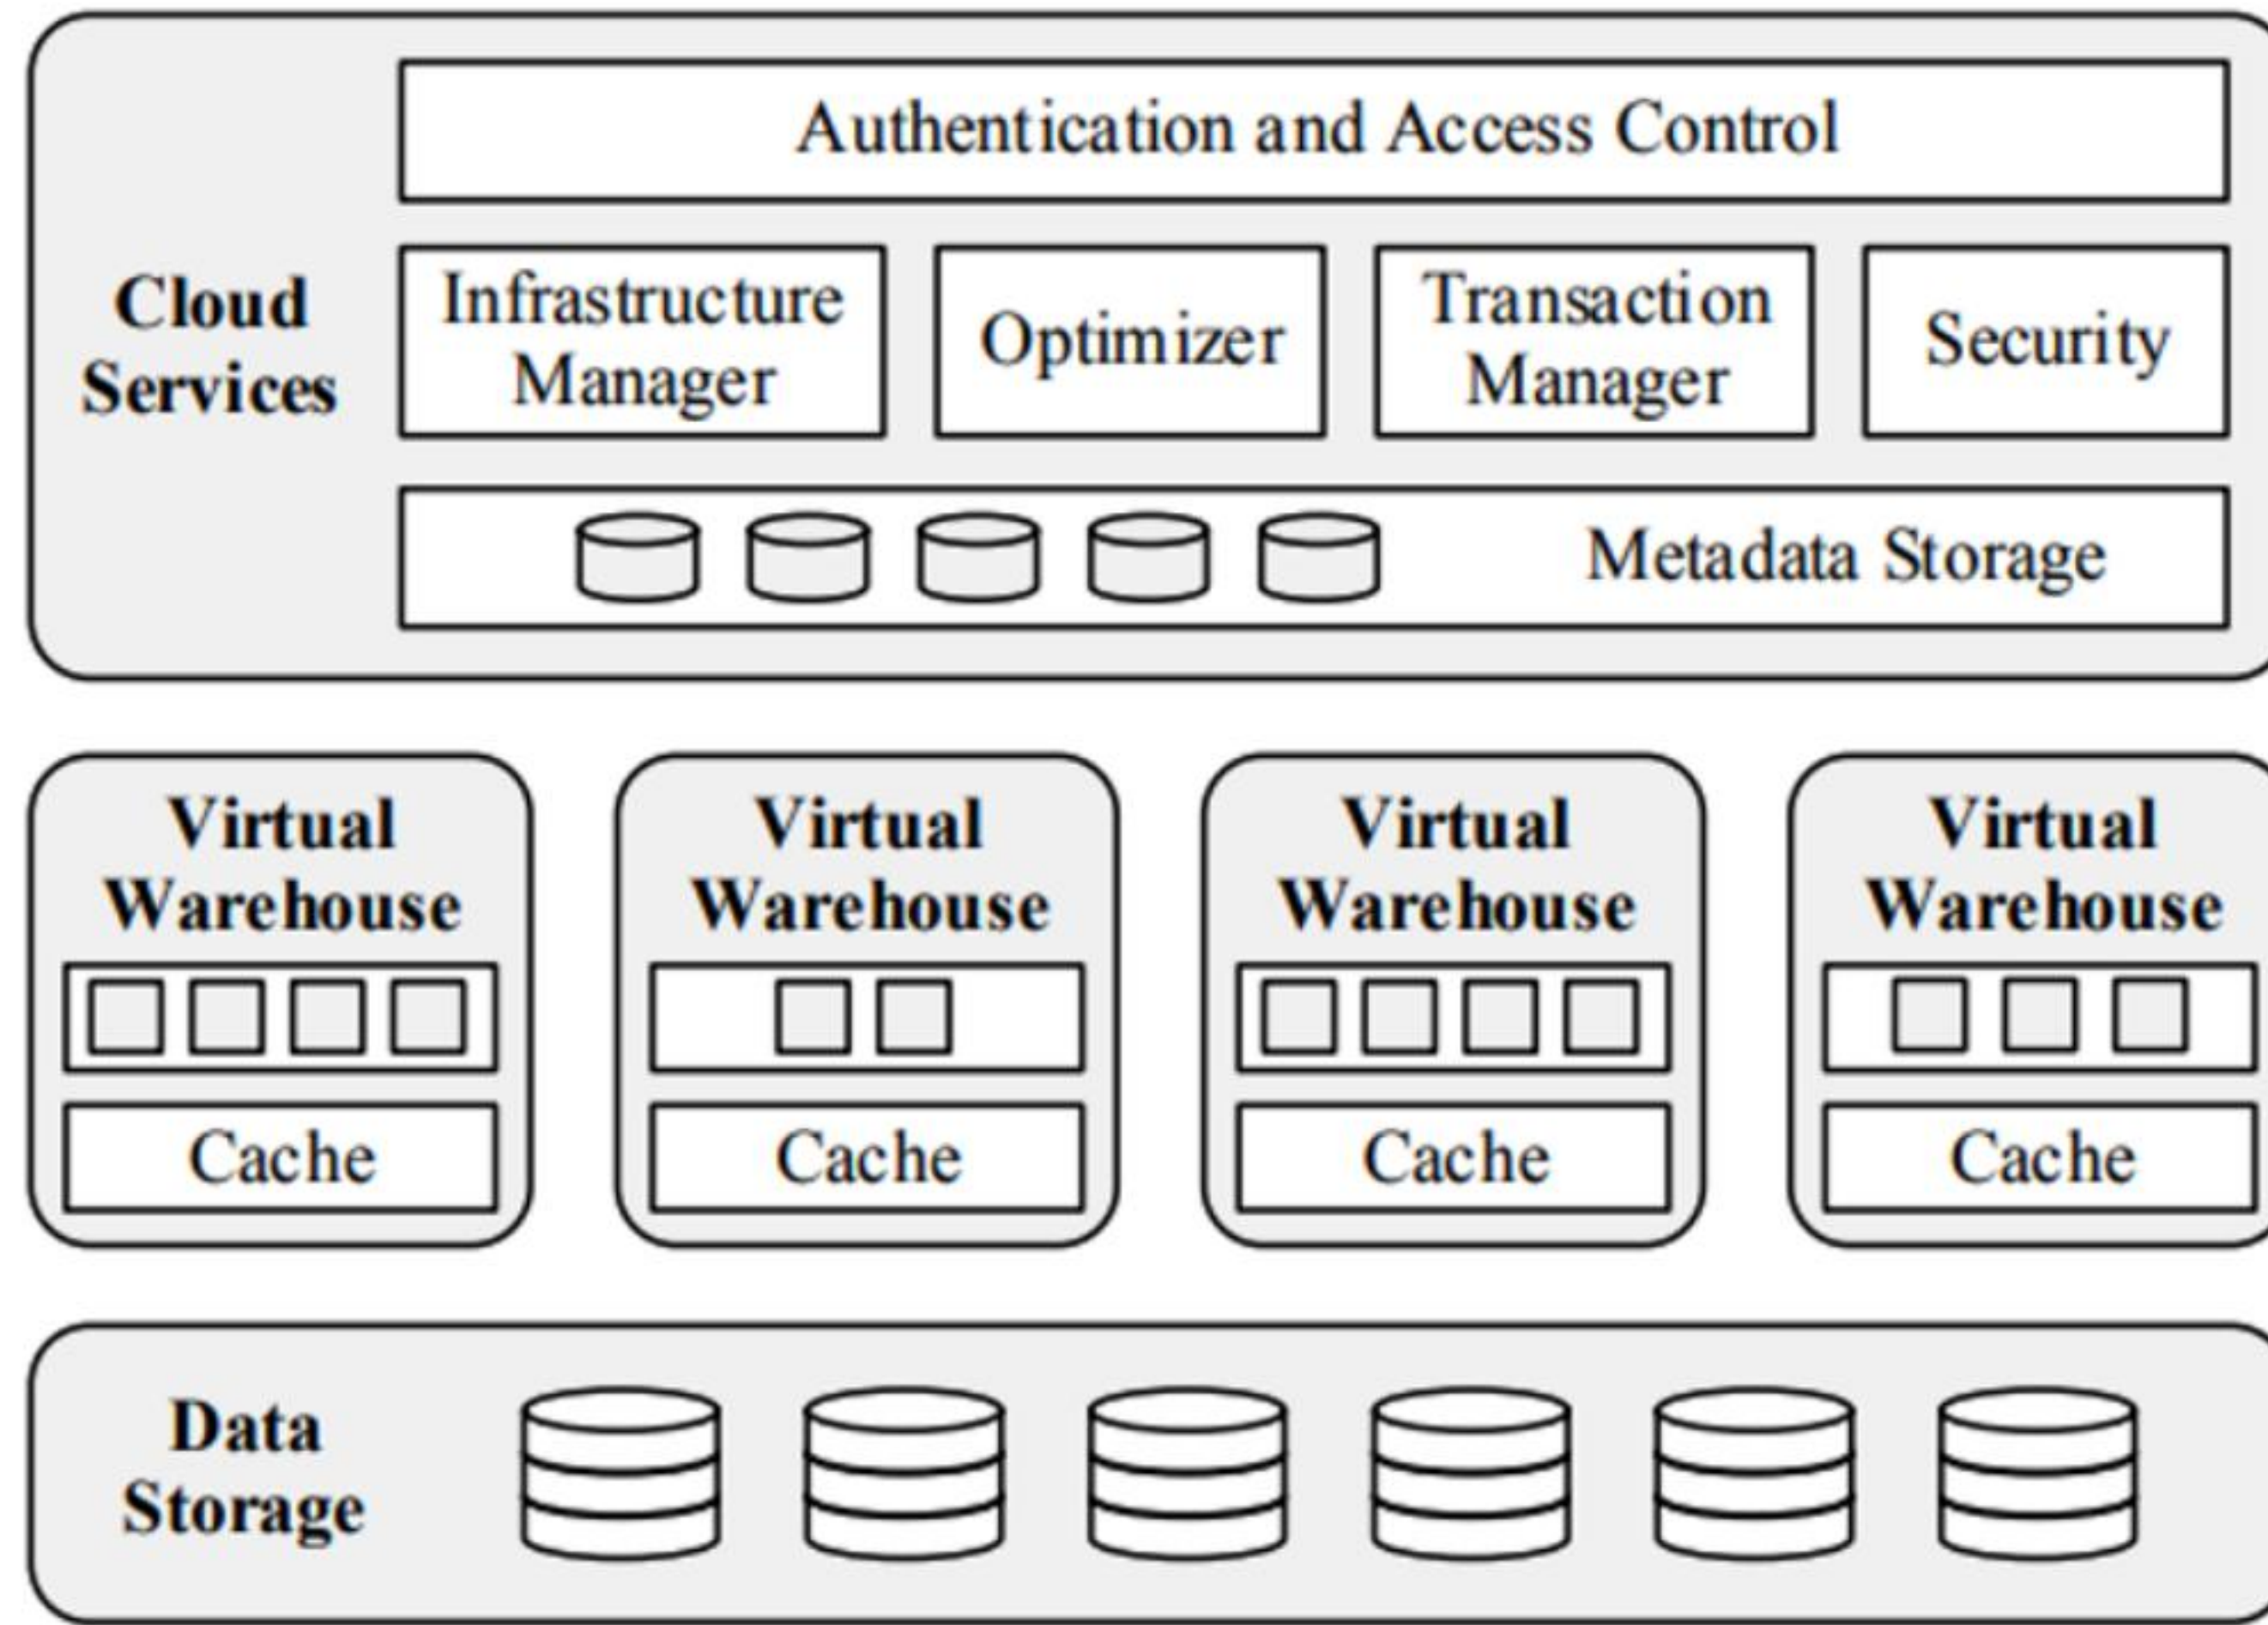
Secure Egress In Serverless Compute

Haowei Yu

Agenda

- Snowflake Introduction
- Serverless in Snowflake
- eBPF Introduction
- How to leverage eBPF to support secure egress
- Closing

Snowflake Introduction - Architecture



Serverless in Snowflake

- Snowflake Abstract Query Execution Resource as Virtual Warehouse
 - “Semi Serverless”
 - Users still need to configure warehouse size (number of servers in cluster)
- Data Loading Service v1
 - Auto-scalable
 - Users only need to define transformation logic (by sql functions)
 - Might co-allocation tasks from different customers on the same physical ec2 instances
 - COGS

Data Loading Service - v3

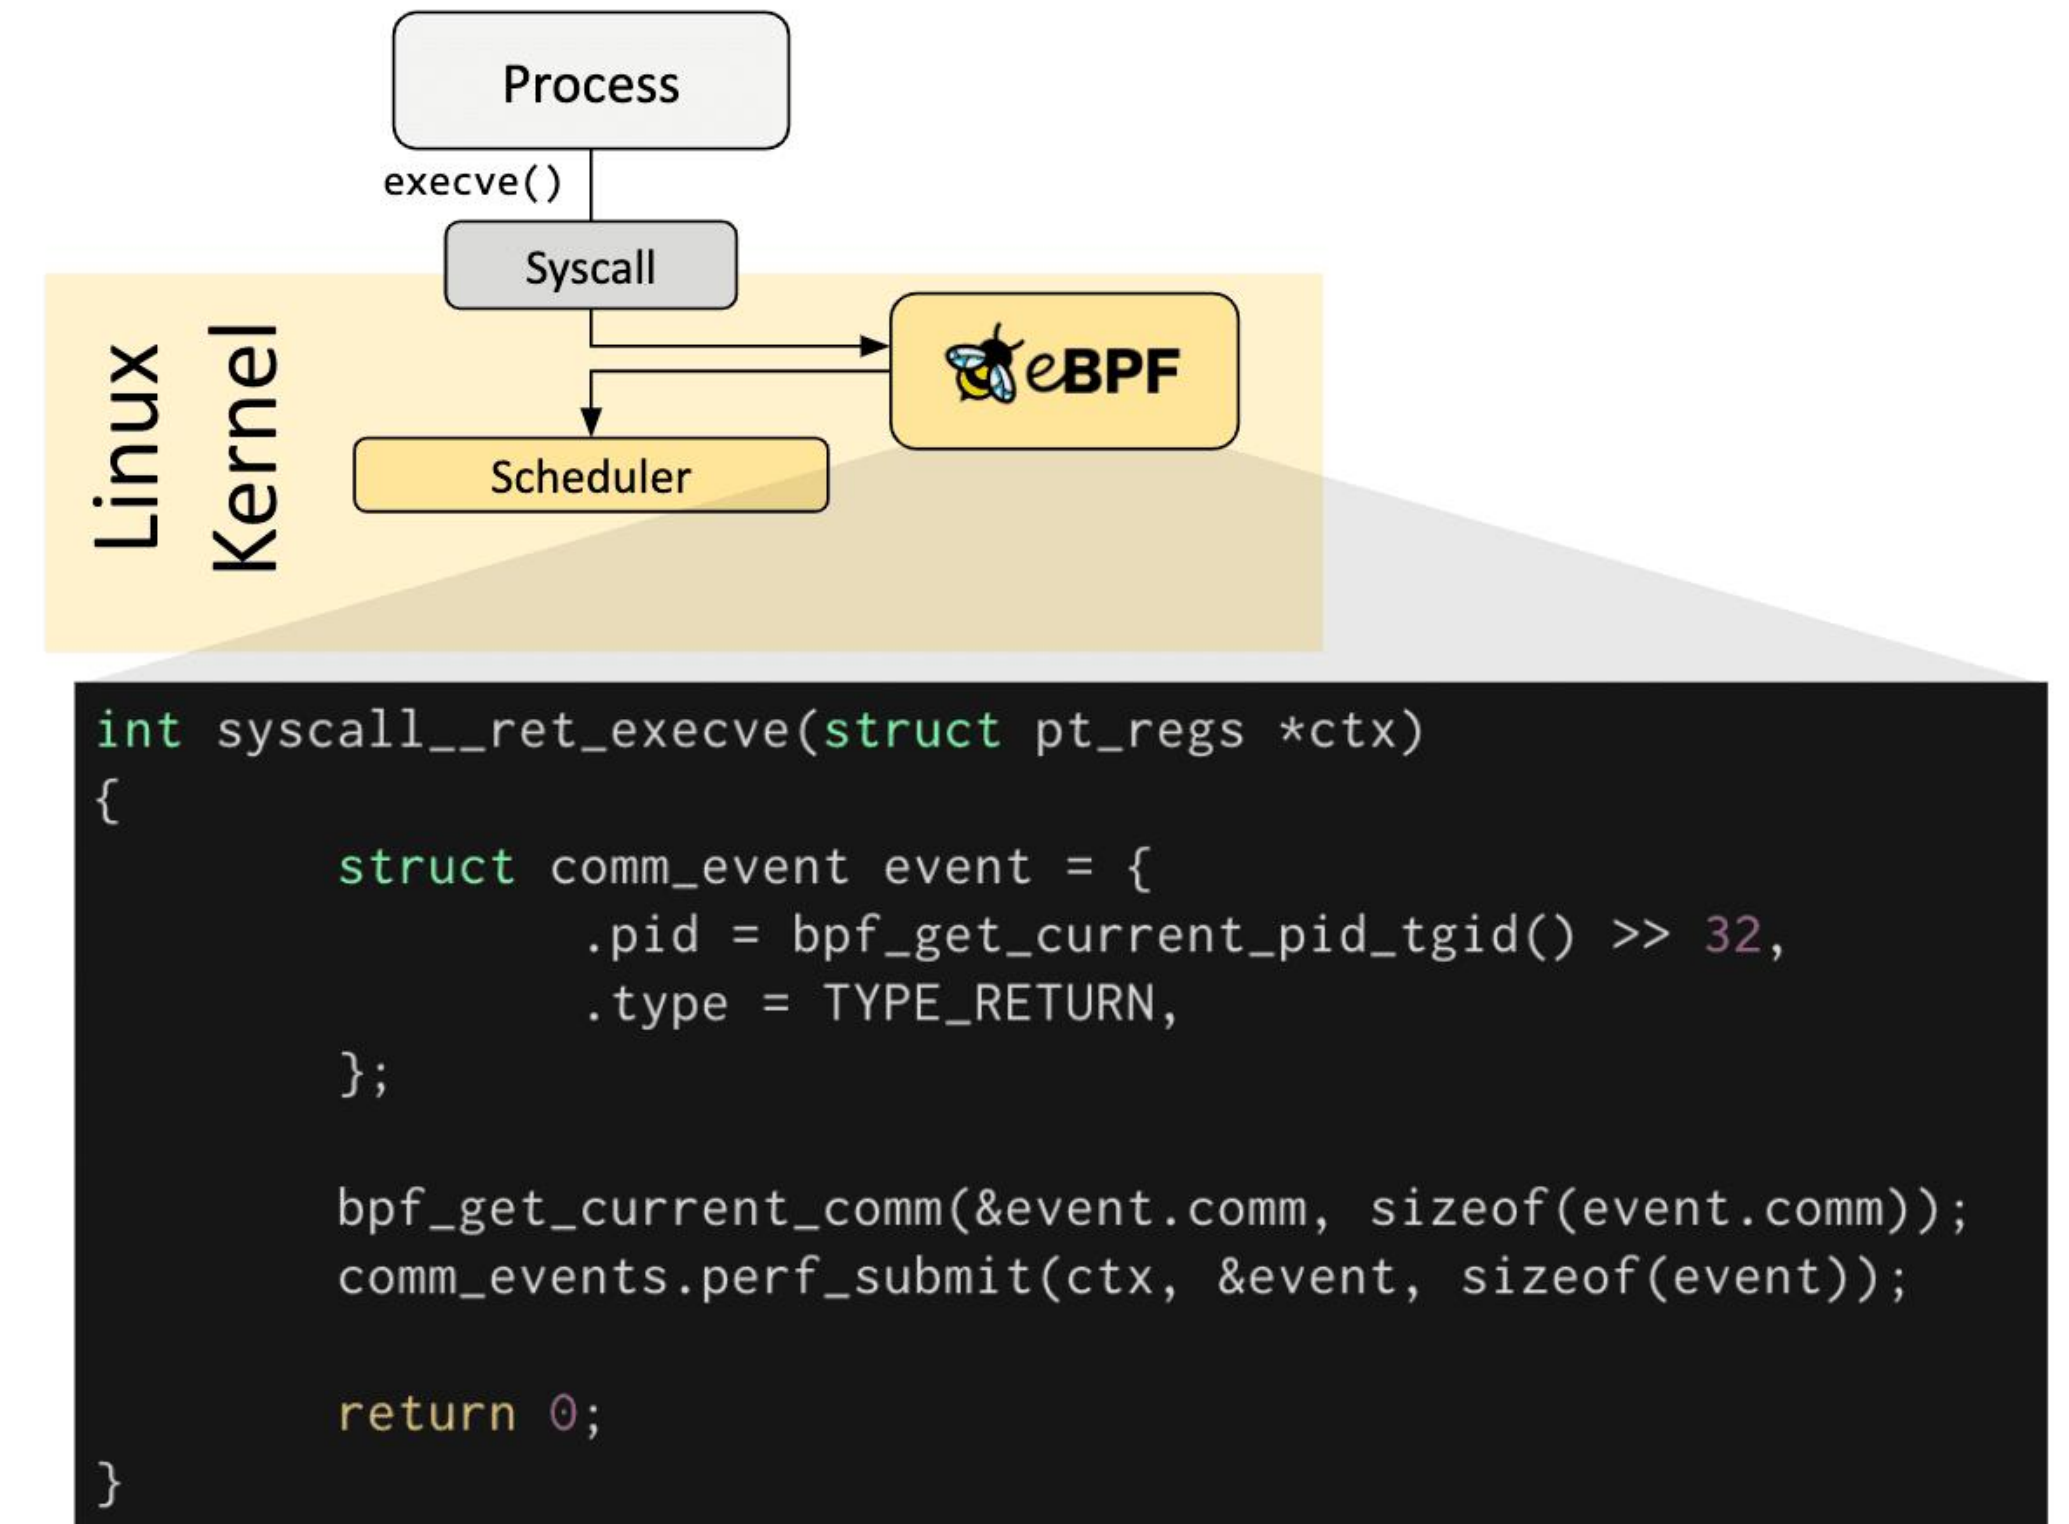
- Users want to define data transformation logic by writing Java/Python Code
- Snowflake expose UDF operator
- Implementation: need an untrusted user code execution environment
- Solution
 - Sandbox user code execution in the local worker node
 - Lock down network access from worker node by default
 - All egress traffic needs to go through a fleet of egress proxies
 - Proxies are also internet gateway
 - Audit
 - Prevent data exfiltration

The problem not solved

- User code needs to connect to endpoint on public internet
 - Geocode
 - Weather Data
 - SaaS Endpoint Connection
- Admin should be able to config which endpoints is (dis)allowed
 - Packet that goes to un-desired endpoints should be dropped
- Proxy is transparent to user code
 - eBPF comes to rescue

eBPF Introduction

- Run sandboxed program in Linux kernel
- Without changing source code
- Event-driven on different hook points
 - System calls
 - Network Event
 - Kprobe



How to use eBPF?

- Write C code and compile with clang
 - Target is eBPF byte code
- Load ebpf byte code with bpf syscall
 - Libbpf wraps different bpf syscall
 - Higher libraries like tc linked with libbpf
- Interact with your eBPF code from userspace with maps
 - Configure the program
 - Collect stats/results
- Verifier
 - Null pointer dereference is not allowed
 - Loop is not allowed in eBPF

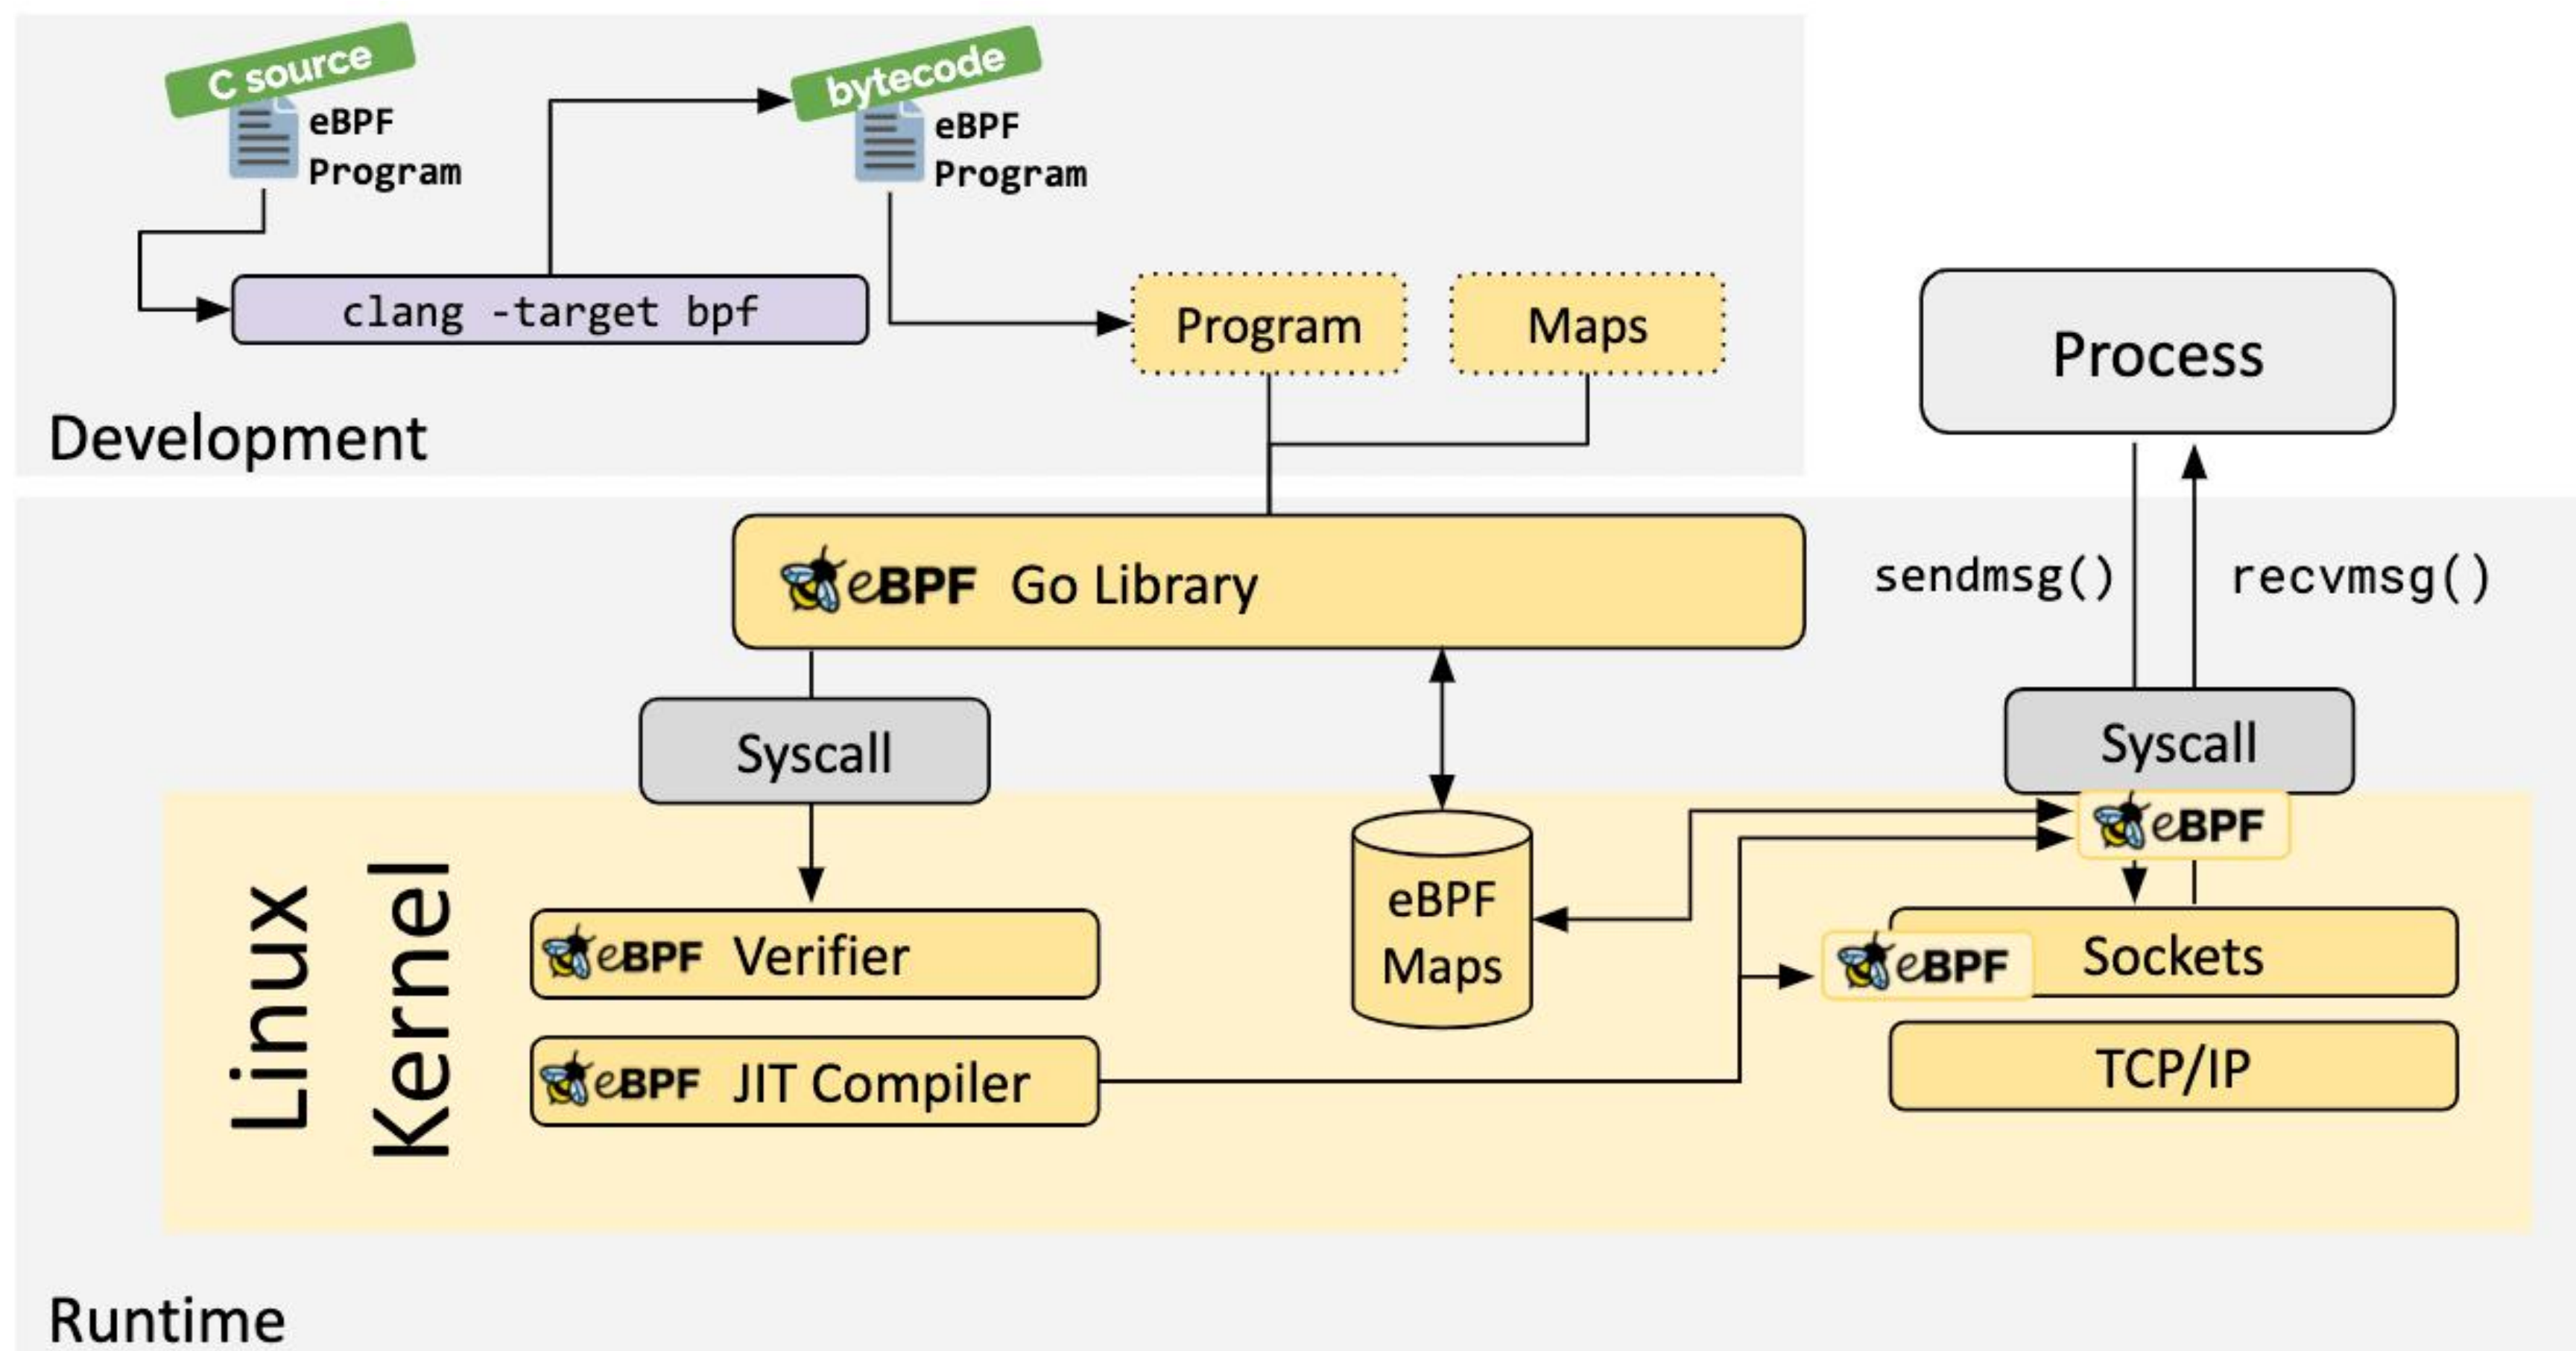
How to use eBPF?

- Write C code and compile with clang
 - Target is eBPF byte code
- Load ebpf byte code with bpf syscall
 - Libbpf wraps different bpf syscall
 - Higher libraries like tc linked with libbpf
- Interact with your eBPF code from userspace with maps
 - Configure the program
 - Collect stats/results
- Verifier
 - Null pointer dereference is not allowed
 - Loop is not allowed in eBPF

eBPF Maps

- Different Data Structures
 - (LRU) Hashmap
 - (LRU) Array
 - LPM
 - Ring Buffer
- Usage
 - A communication channel between kernel and userspace
 - Collect time spent on each syscall
 - Share state between different ebpf function invocation
 - Share state between different packet processing

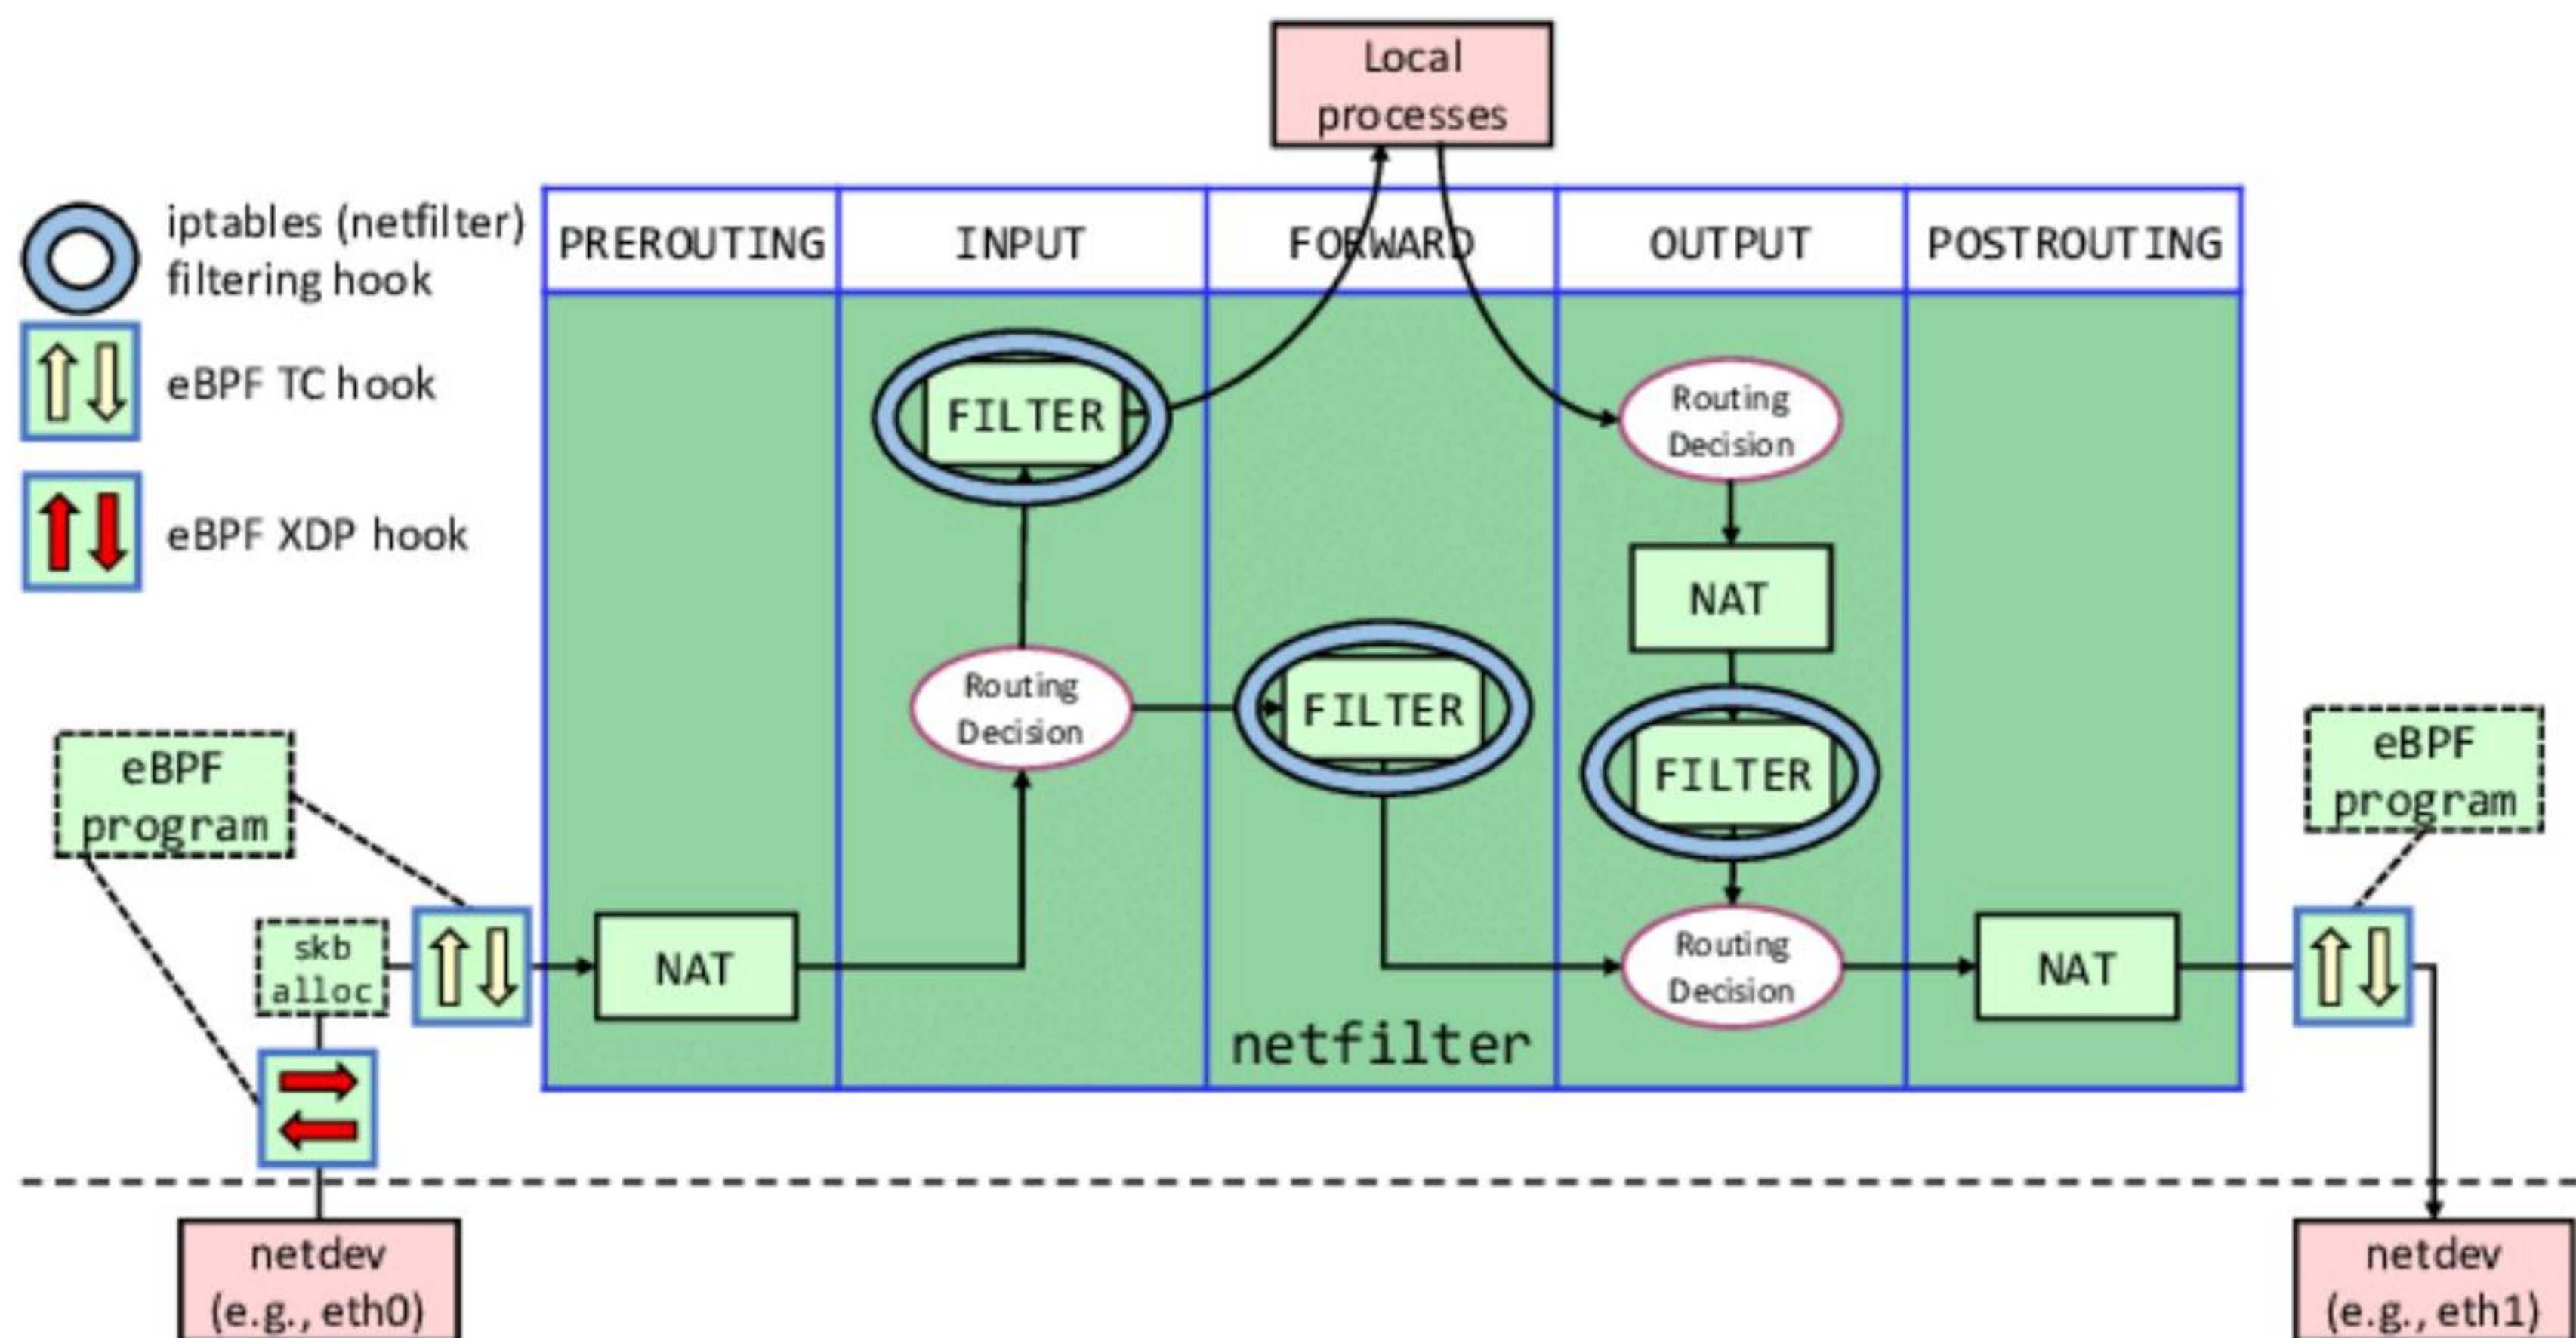
Put it all together



eBPF In Networking

- Monitoring
 - Within network stack on a single machine
 - Between machines inside a cluster
- Packet Forwarding/Redirect
 - Easier to manager than IP table
- Rate limiting
- Policy Enforcement

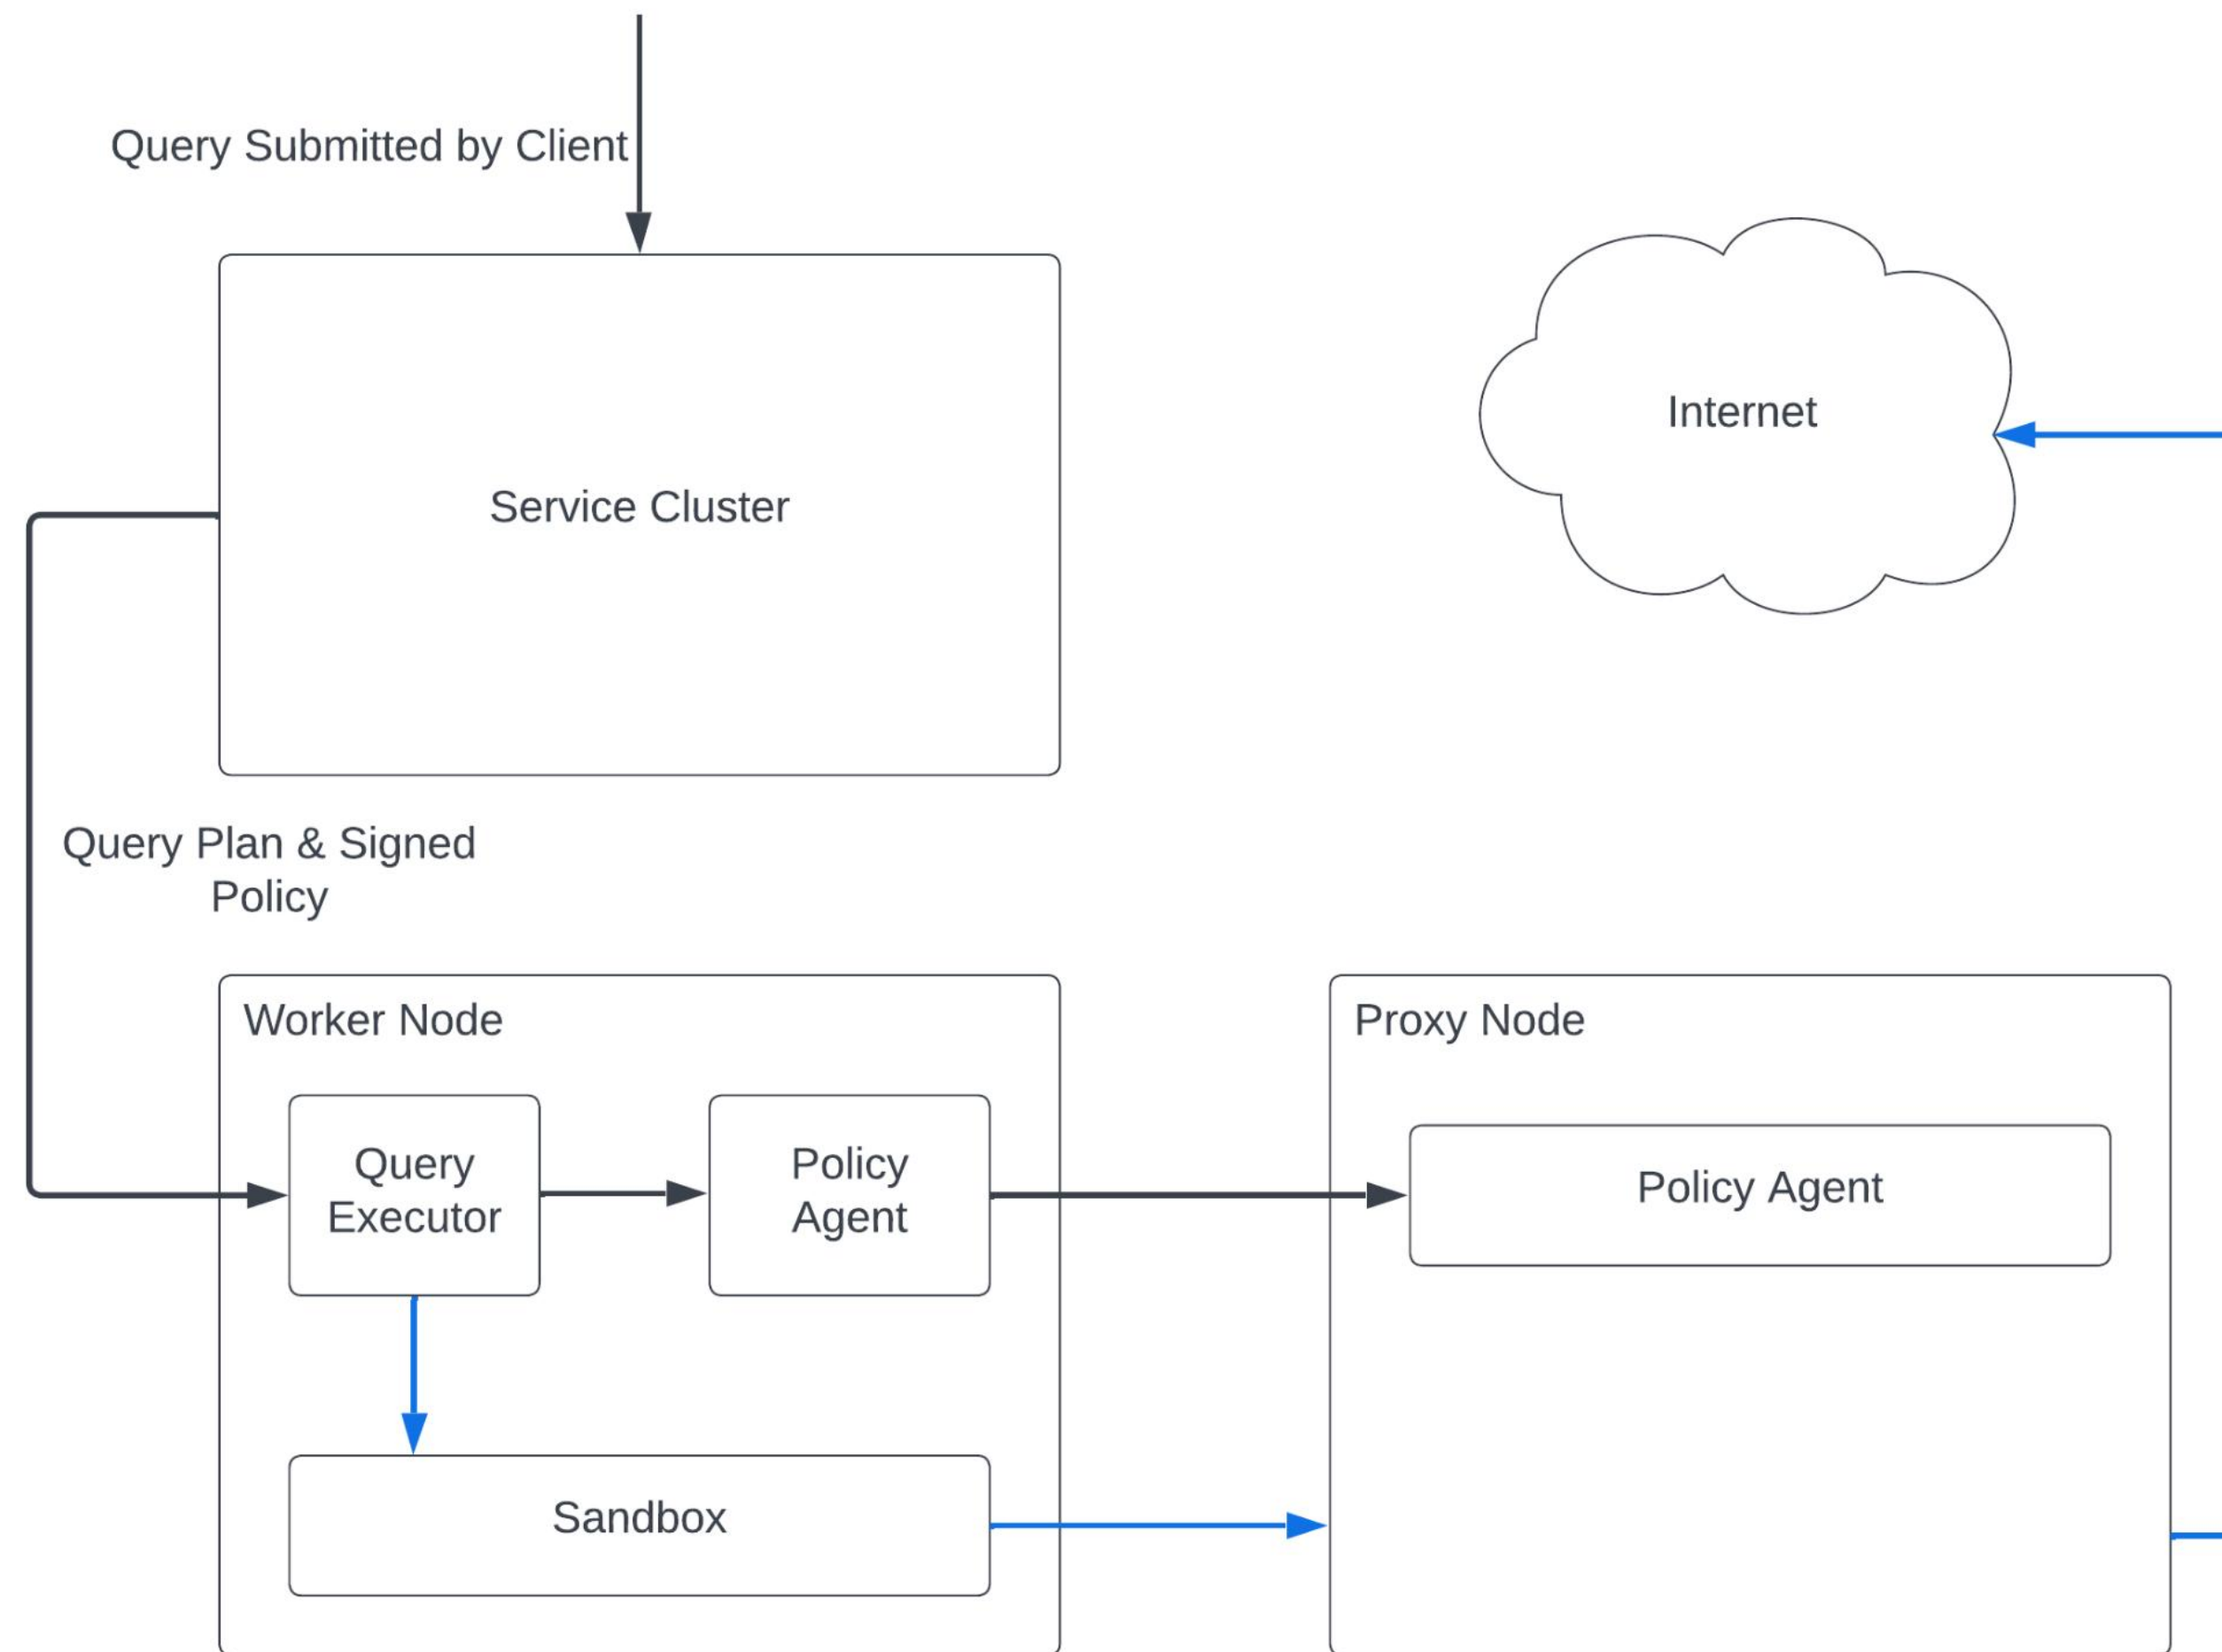
Attach eBPF in XDP and TC



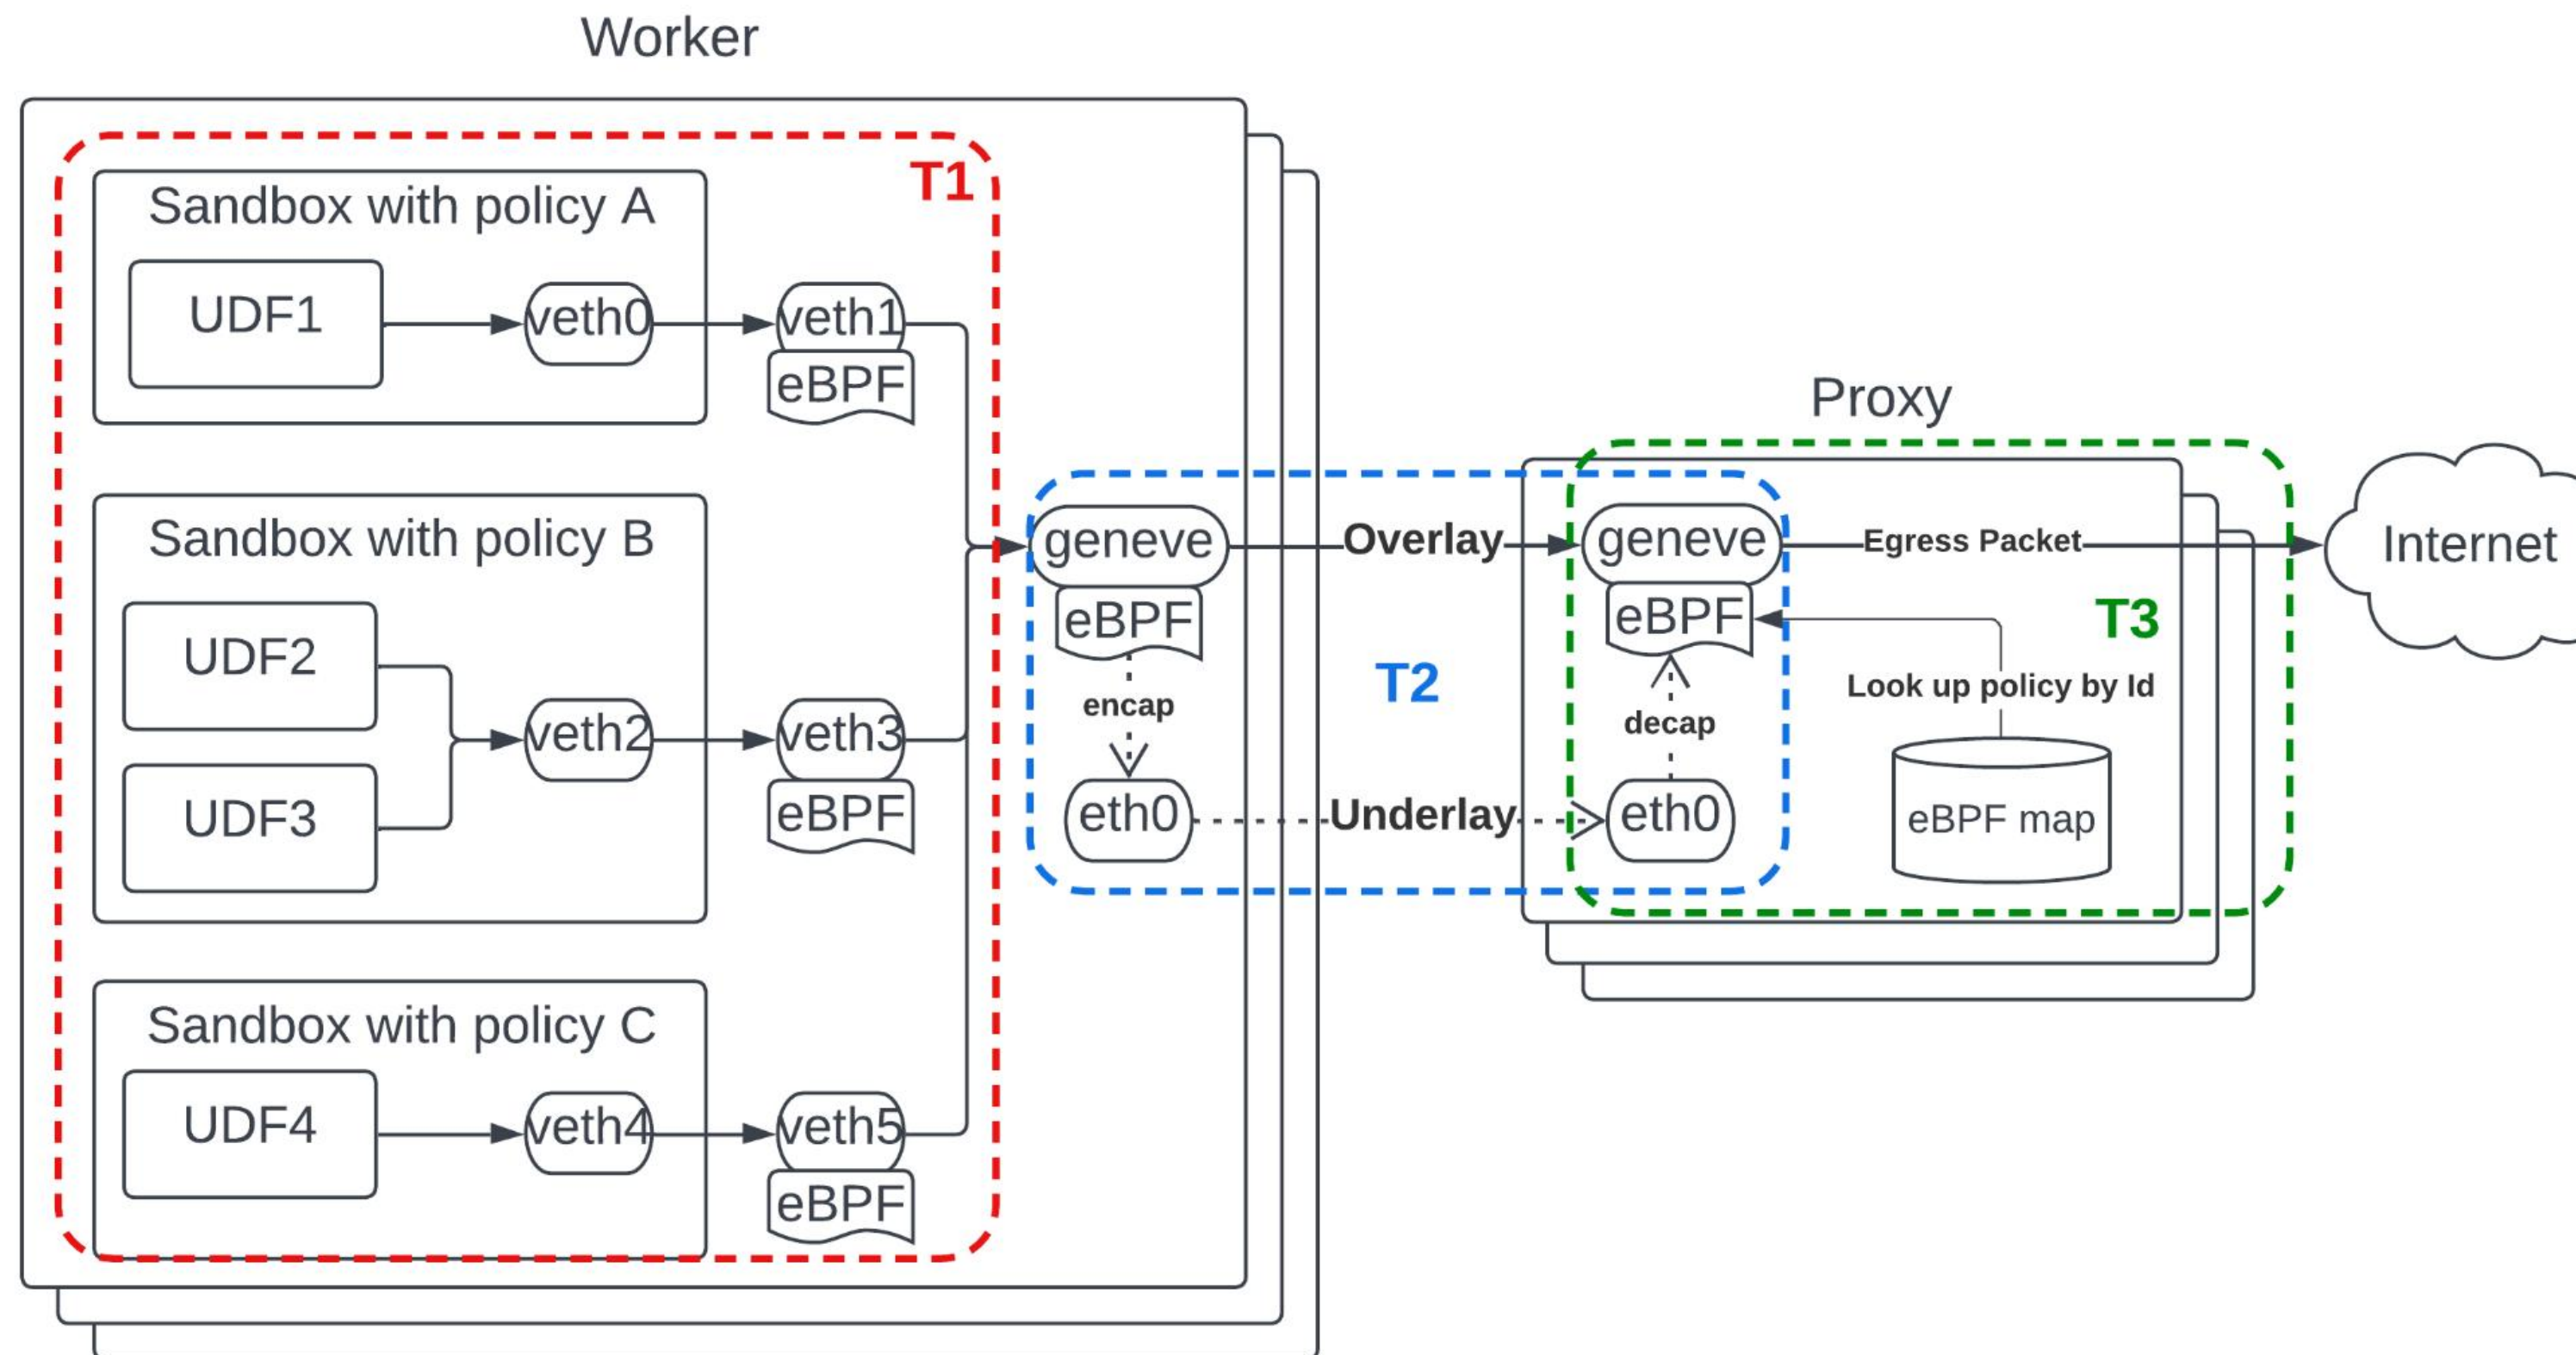
eBPF In Networking

- Monitoring
 - Within network stack on a single machine
 - Between machines inside a cluster
- Packet Forwarding/Redirect
 - Easier to manager than IP table
- Rate limiting
- Policy Enforcement

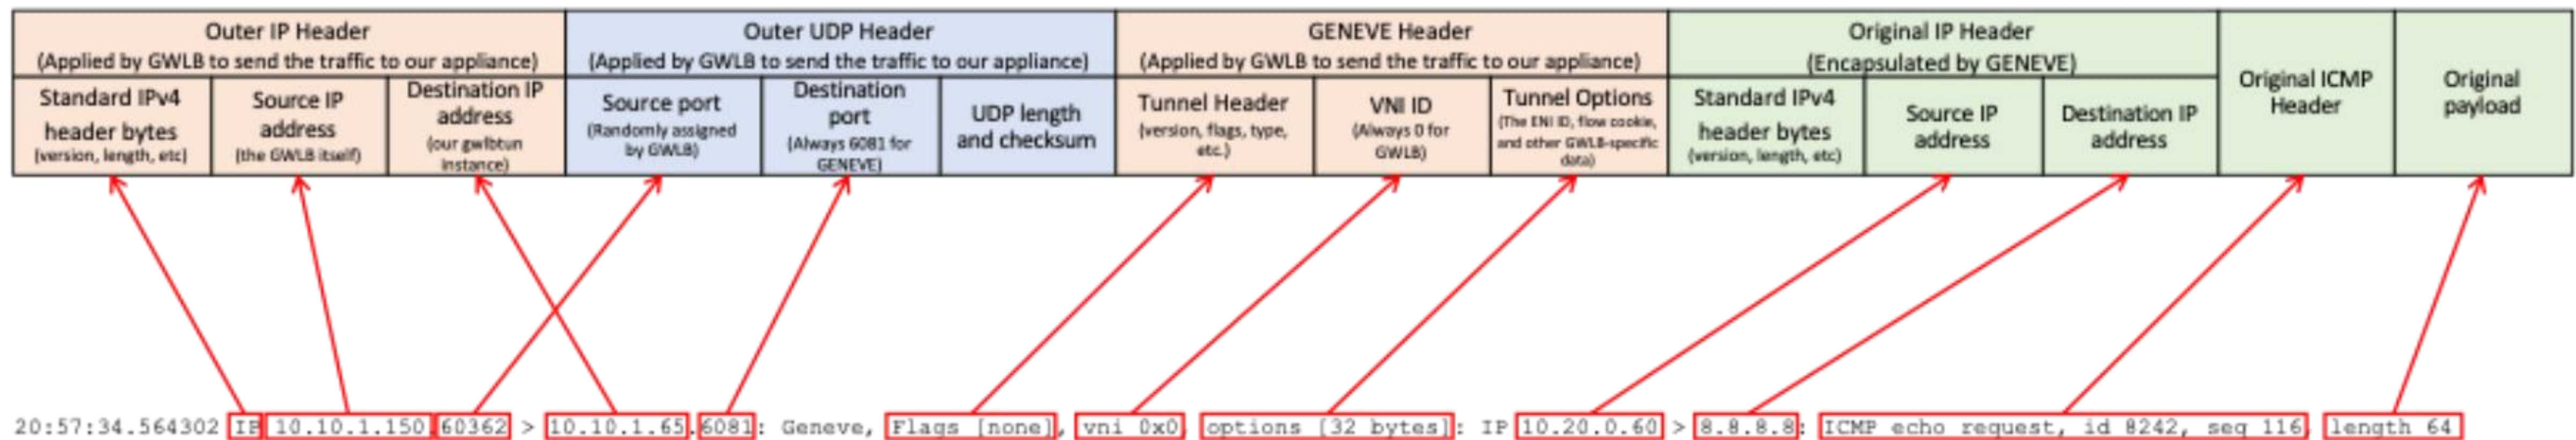
Solution - Control Flow



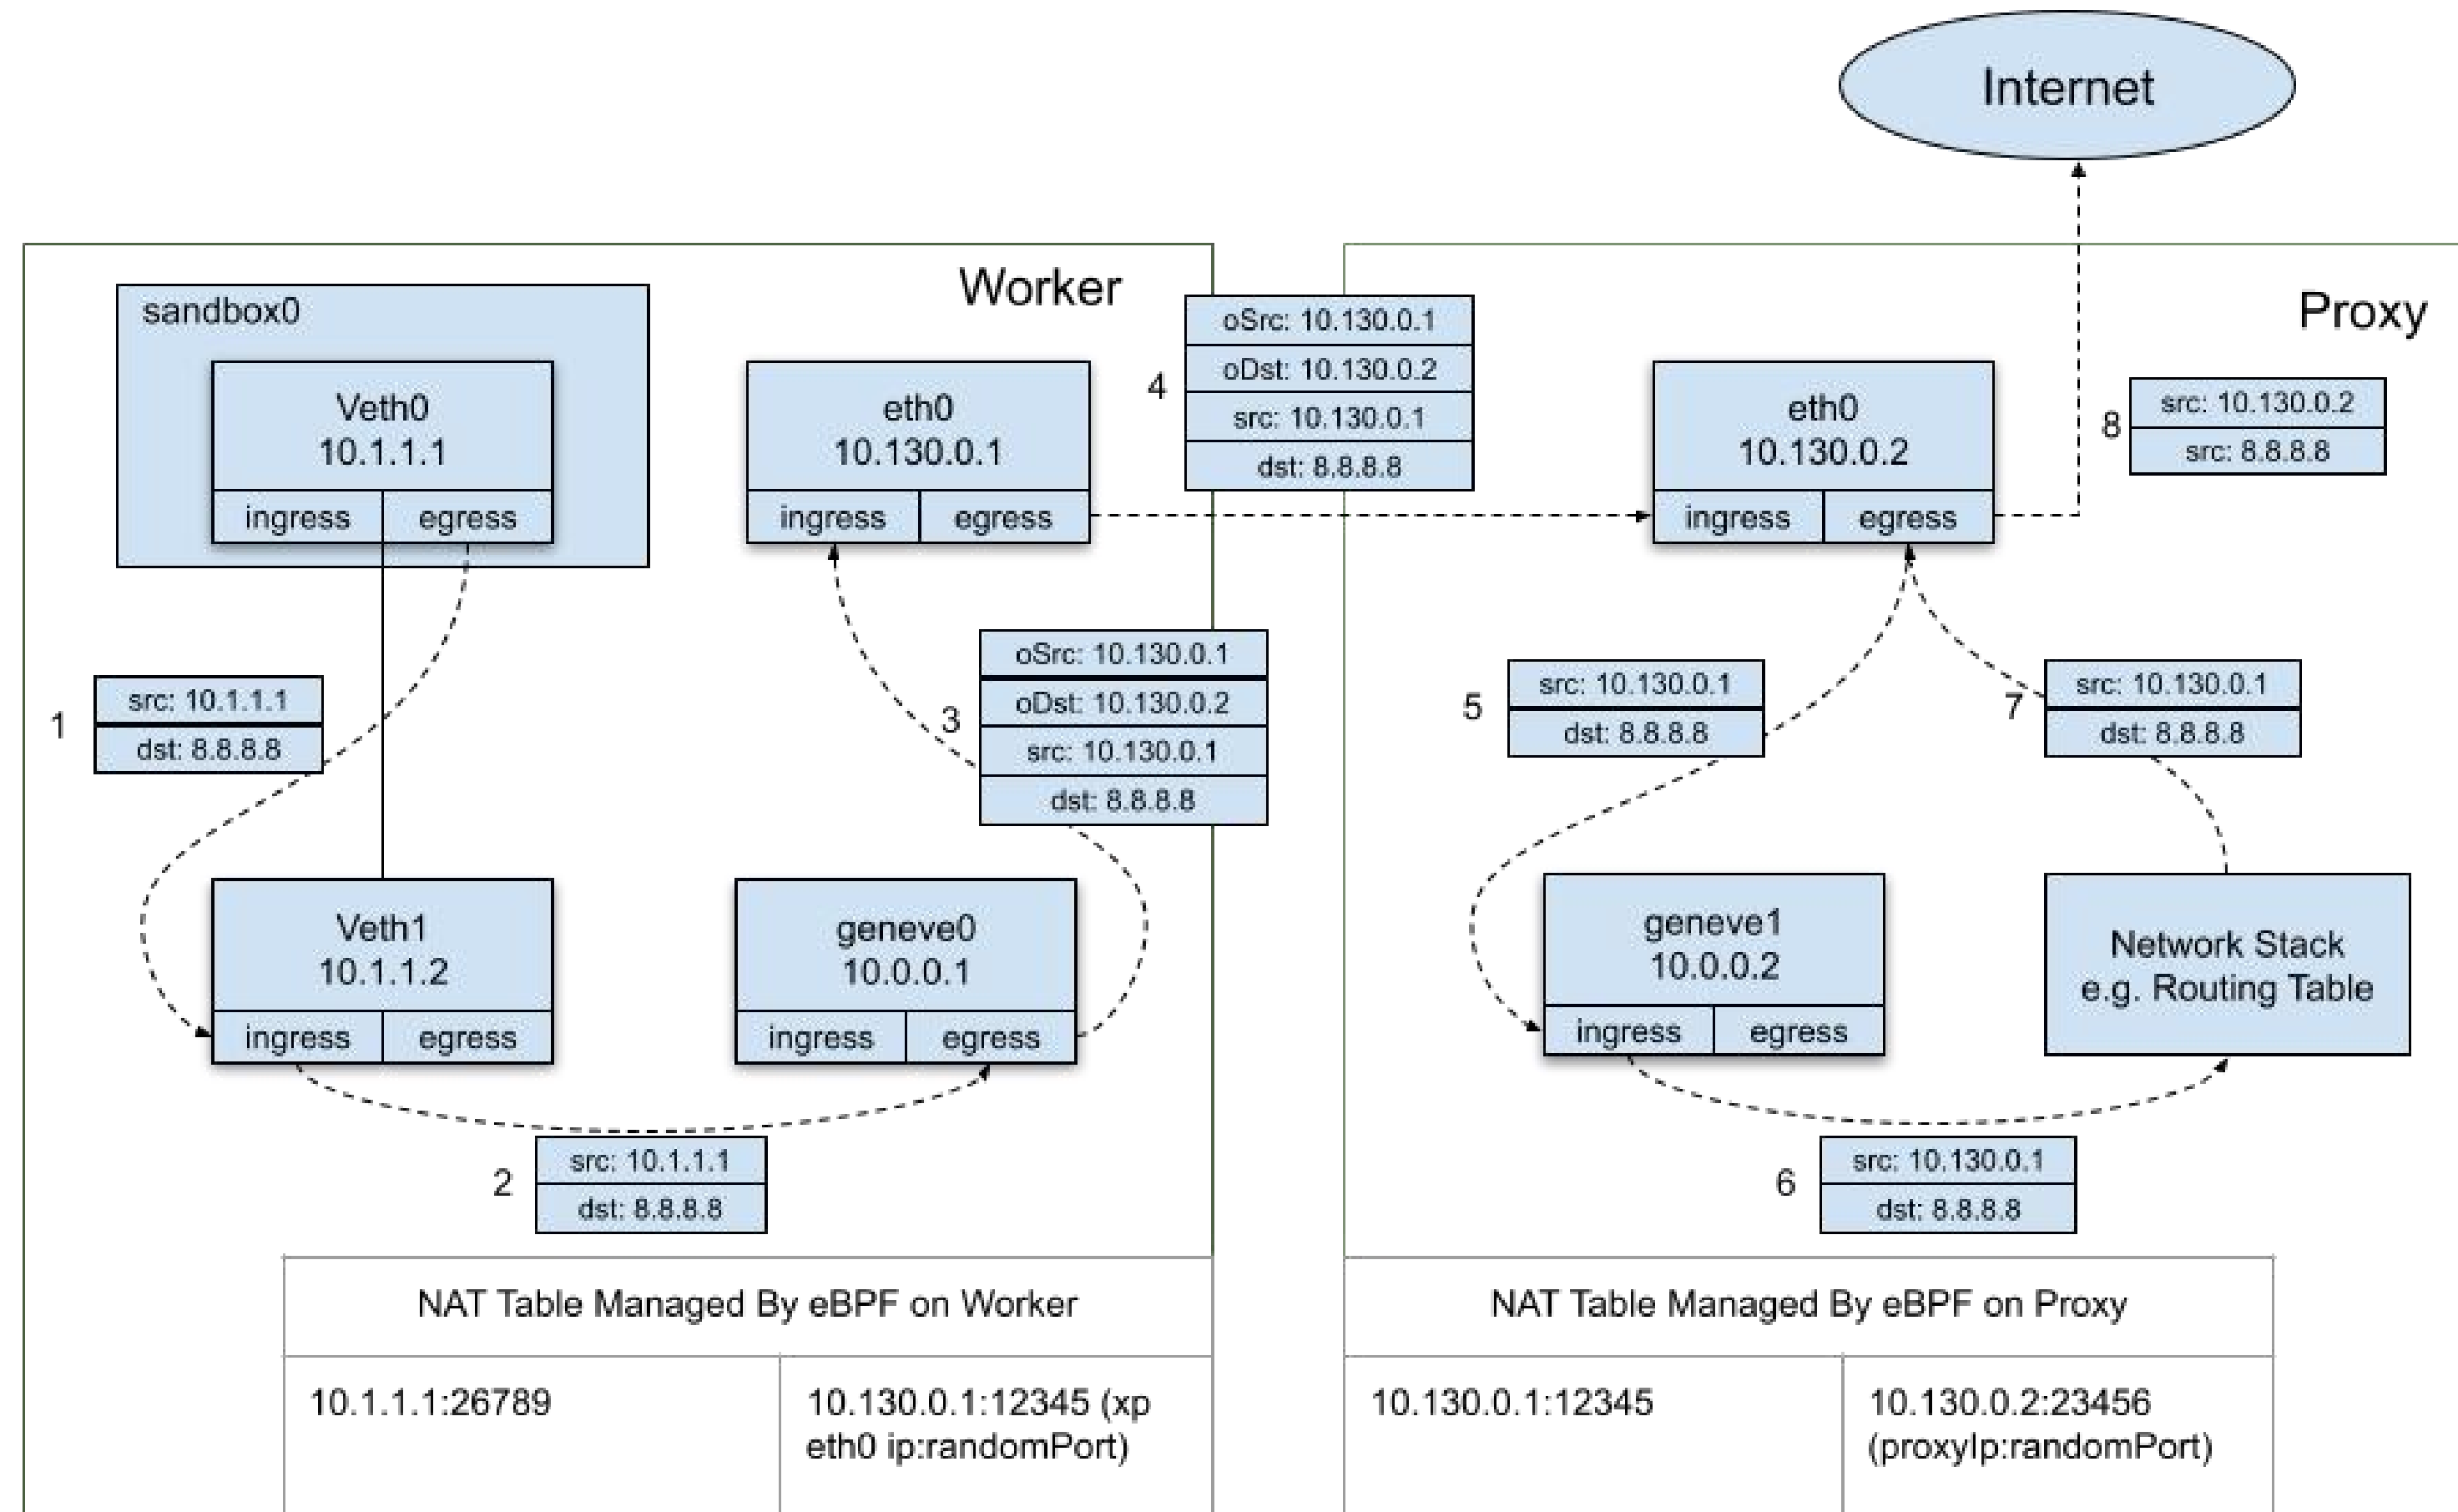
Solution - High Level



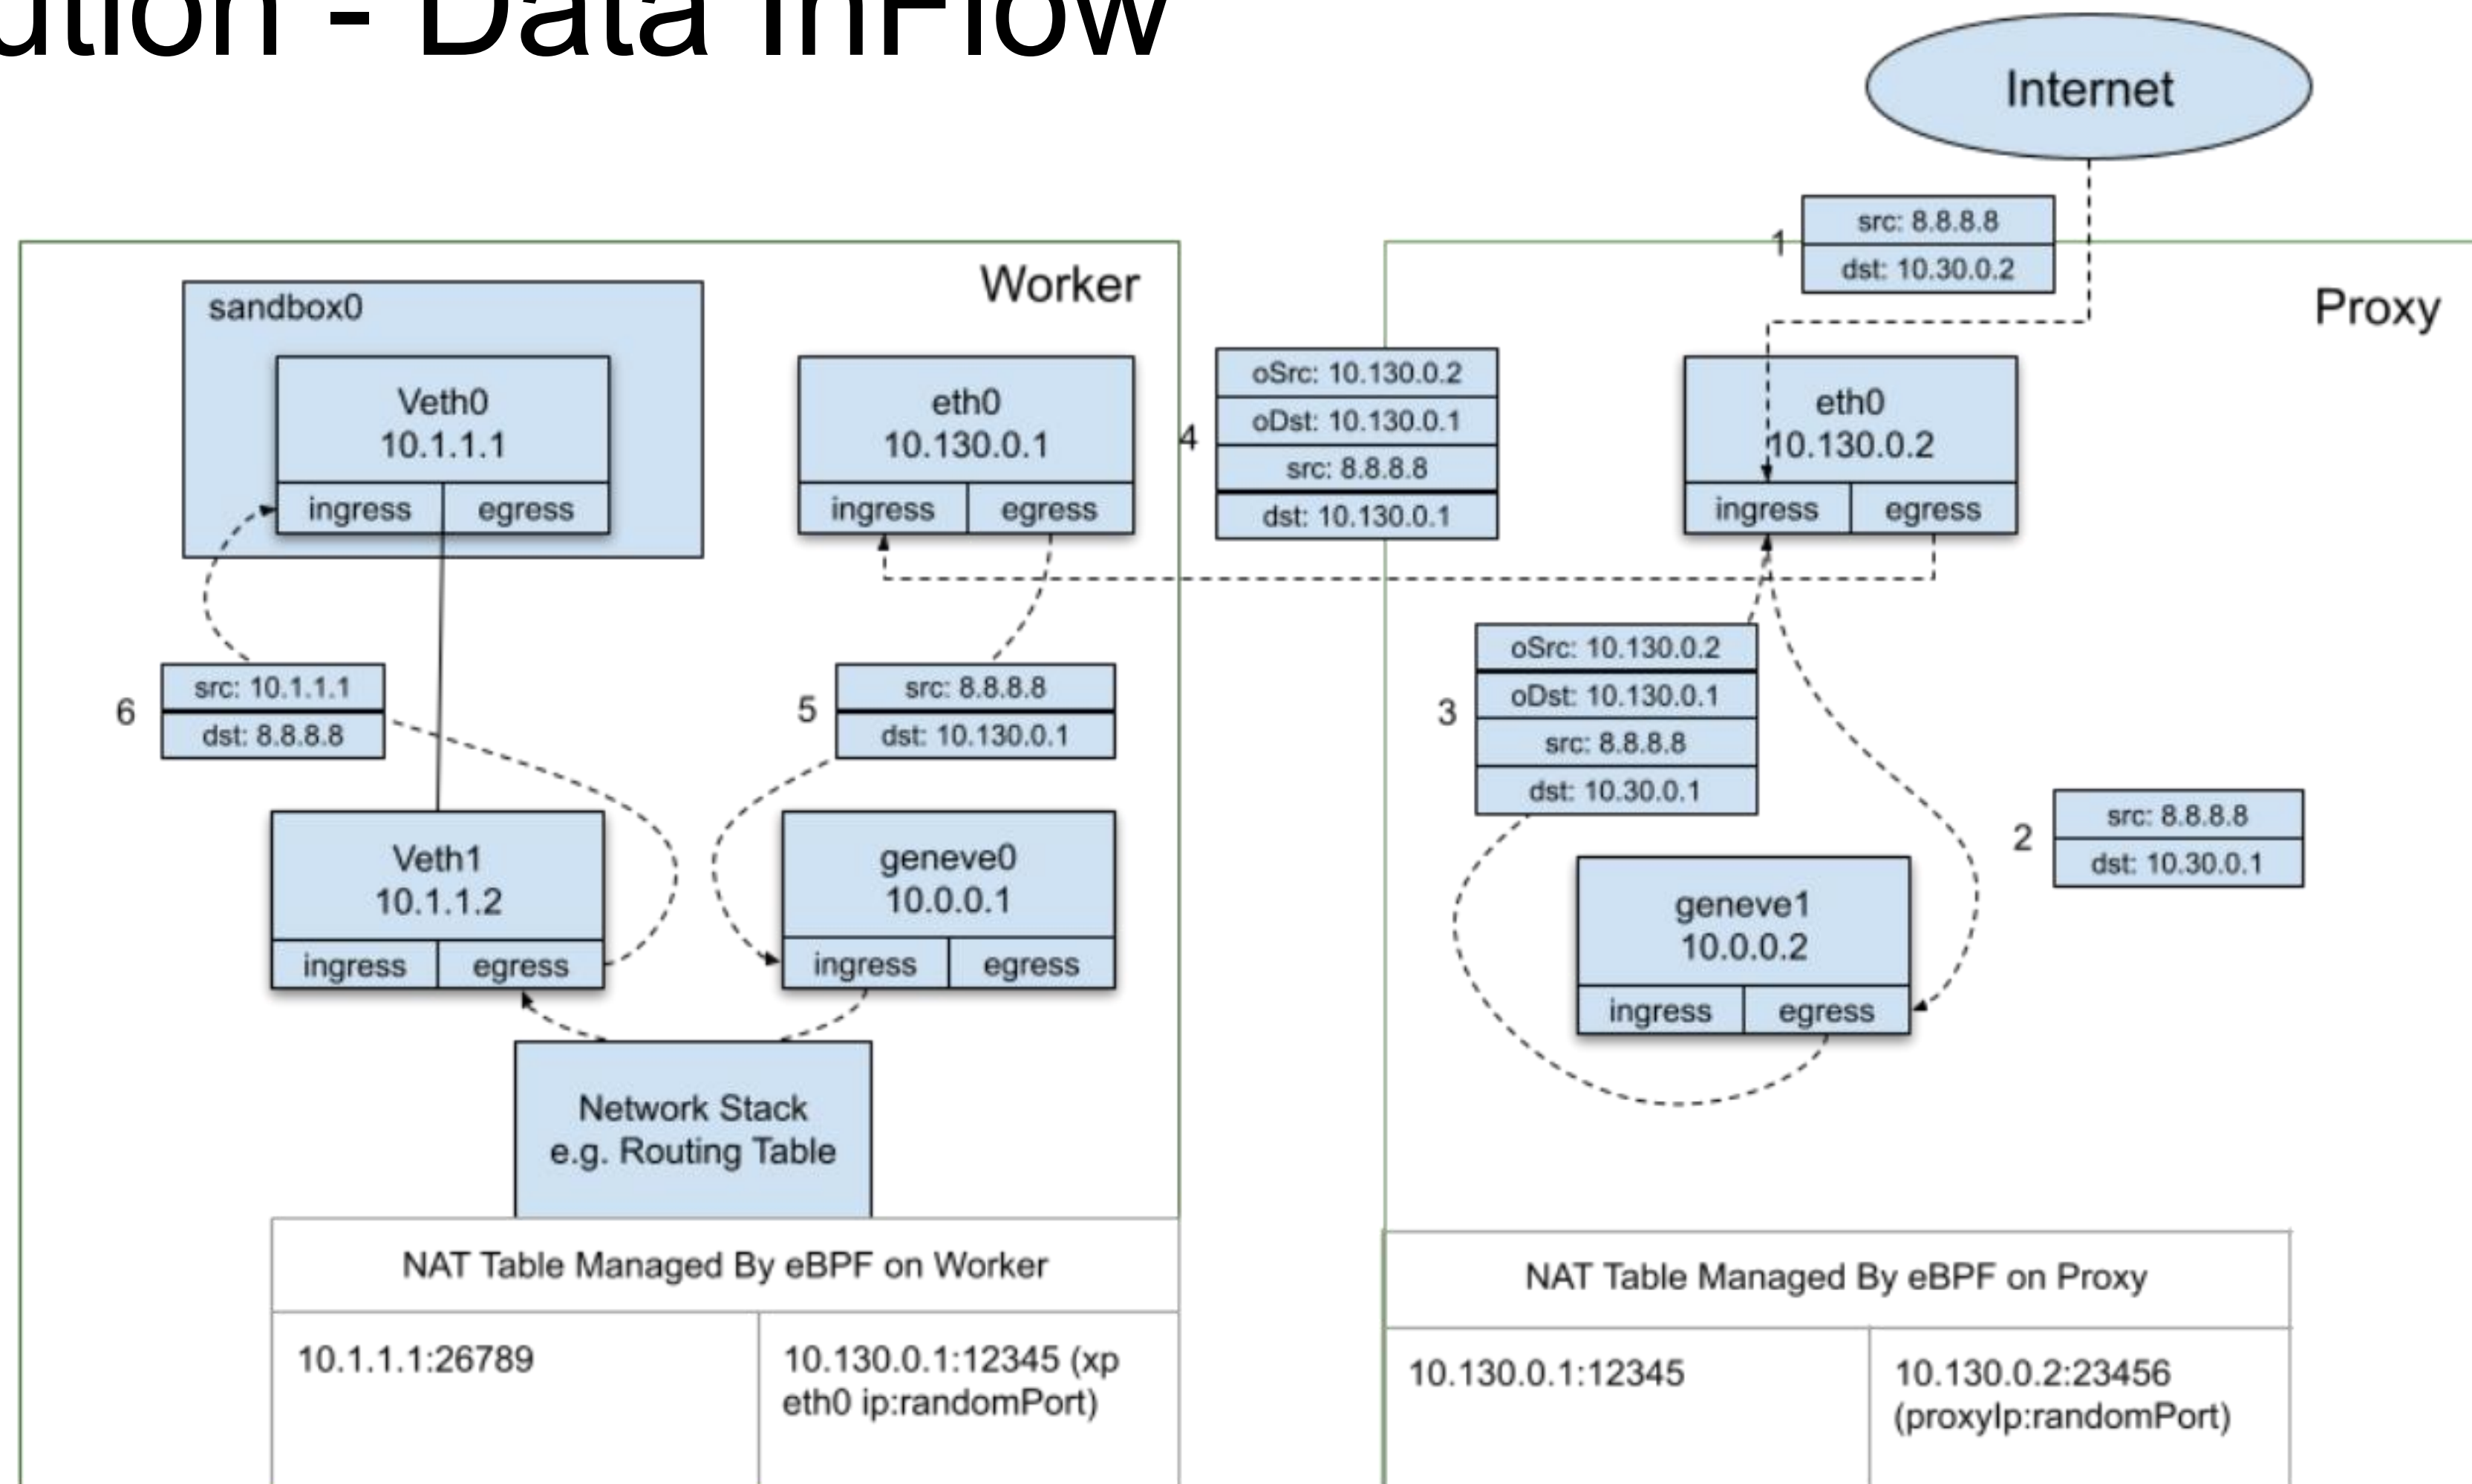
Solution - Packet Encapsulation



Solution - Data OutFlow



Solution - Data InFlow



Closing

- Connection to private endpoint
- Other stats/metric that can be collected from eBPF
- LSM BPF for advanced security features

THANKS

软件正在重新定义世界

Software Is Redefining The World